

GS Operations - C10M

Concurrent 10 million connection and Load Test - DPDK/mTCP/Pktgen

Project Name	C10M
Purpose	<p>Develop an opensource based packet generator for BIGIP load evaluation/trouble-shooting using technology below</p> <p>DPDK http://www.dpdk.org/ is an user space drivers and libraries for fast packet processing, it can generates 10M/pps, 10M/cps</p> <p>mTCP A Highly Scalable User-level TCP Stack for Multicore Systems.</p> <p>MoonGen https://github.com/emmericp/MoonGen to generate raw packet like SYN/RST/ACK/UDP/ICMP flooding</p> <p>Pktgen-DPDK http://dpdk.org/browse/apps/pktgen-dpdk/ Another raw packet generator from Intel, runs with latest DPDK repo</p>
Project Contacts	Vincent Li

Background:

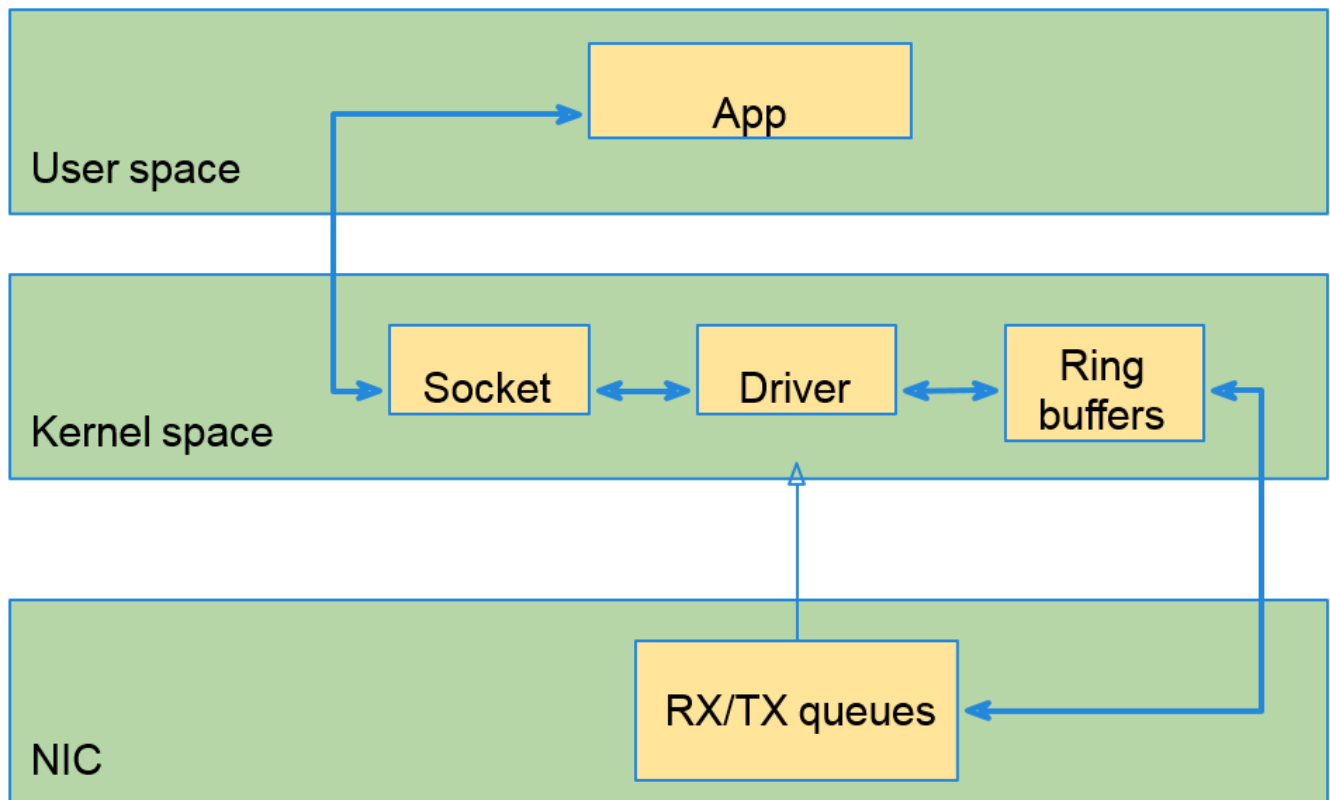
The Secret to 10 Million Concurrent Connections -The Kernel is the Problem, Not the Solution

<http://highscalability.com/blog/2013/5/13/the-secret-to-10-million-concurrent-connections-the-kernel-i.html>

Problem:

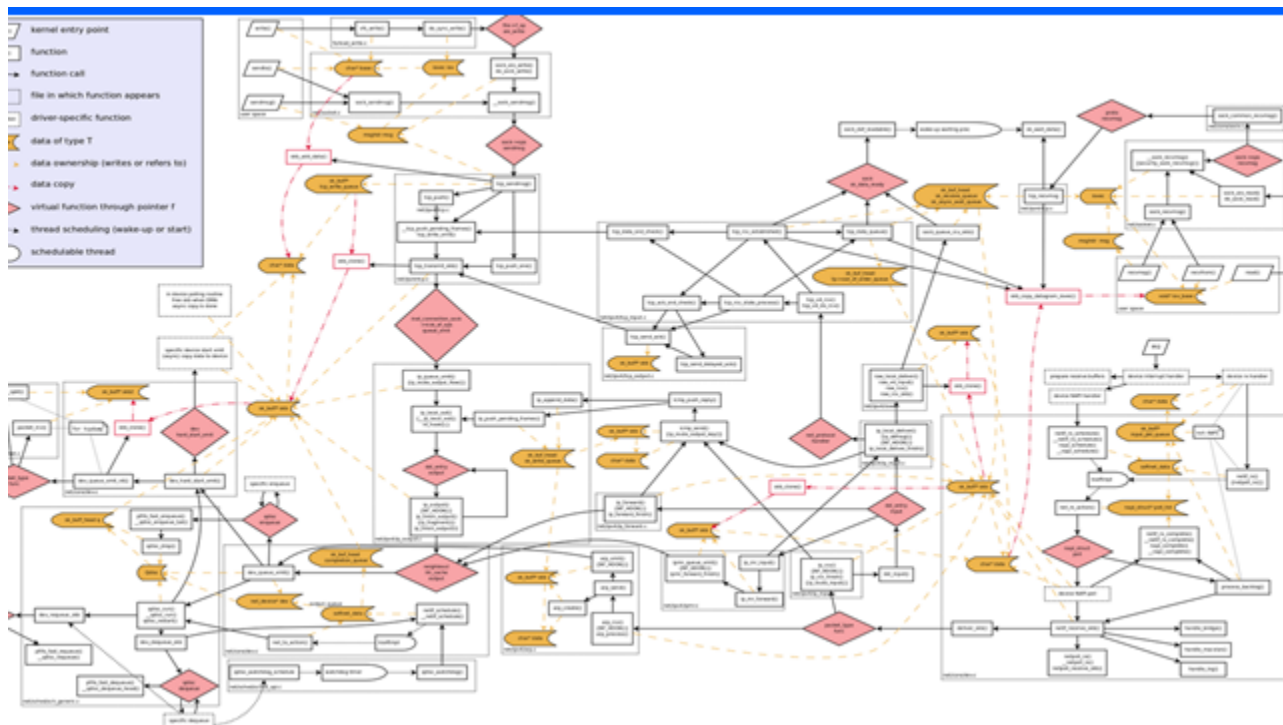
http://www.slideshare.net/garyachy/dpdk-44585840?qid=254b419f-1d44-44f1-99c4-87f13b7d5fe4&v=default&b=&from_search=8

Simplified packet processing in Linux:



Real packet processing in Linux:

Linux network data flow



- System calls
- Context switching on blocking I/O
- Data Copying from kernel to user space
- Interrupt handing in kernel

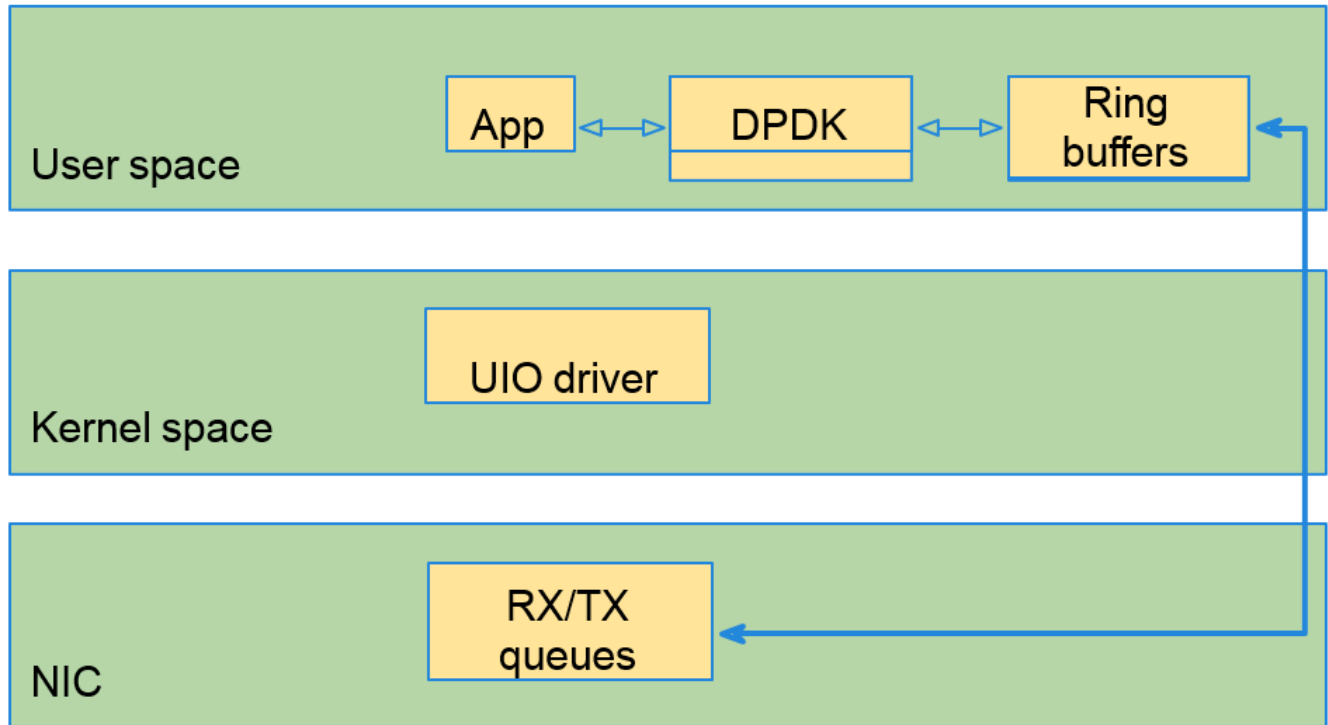
Expense of sendto :

- sendto - system call: 96ns
- sosend_dgram - lock sock_buff, alloc mbuf, copy in: 137ns
- udp_output - UDP header setup: 57ns
- ip_output - route lookup, ip header setup: 198ns
- ether_output - MAC lookup, MAC header setup: 162ns
- ixgbe_xmit - device programming: 220ns

Total: 950ns

Solution:

Packet processing with DPDK



- Processor affinity (separate cores)
- Huge pages(no swap, TLB)
- UIO (no copying from kernel)
- Polling (no interrupts overhead)
- Lockless synchronization(avoid waiting)
- Batch packets handling
- SSE, NUMA awareness

UIO for example:

Kernel space (UIO framework) <----->/dev/uioX<----->userspace epoll/mmap<----->App

Problem:

<http://www.ndsl.kaist.edu/~kyoungsoo/papers/mtcp.pdf>

Limitations of the Kernel's TCP stack

- Lack of connection locality
- Shared file descriptor space
- Inefficient per-packet processing

- System call overhead

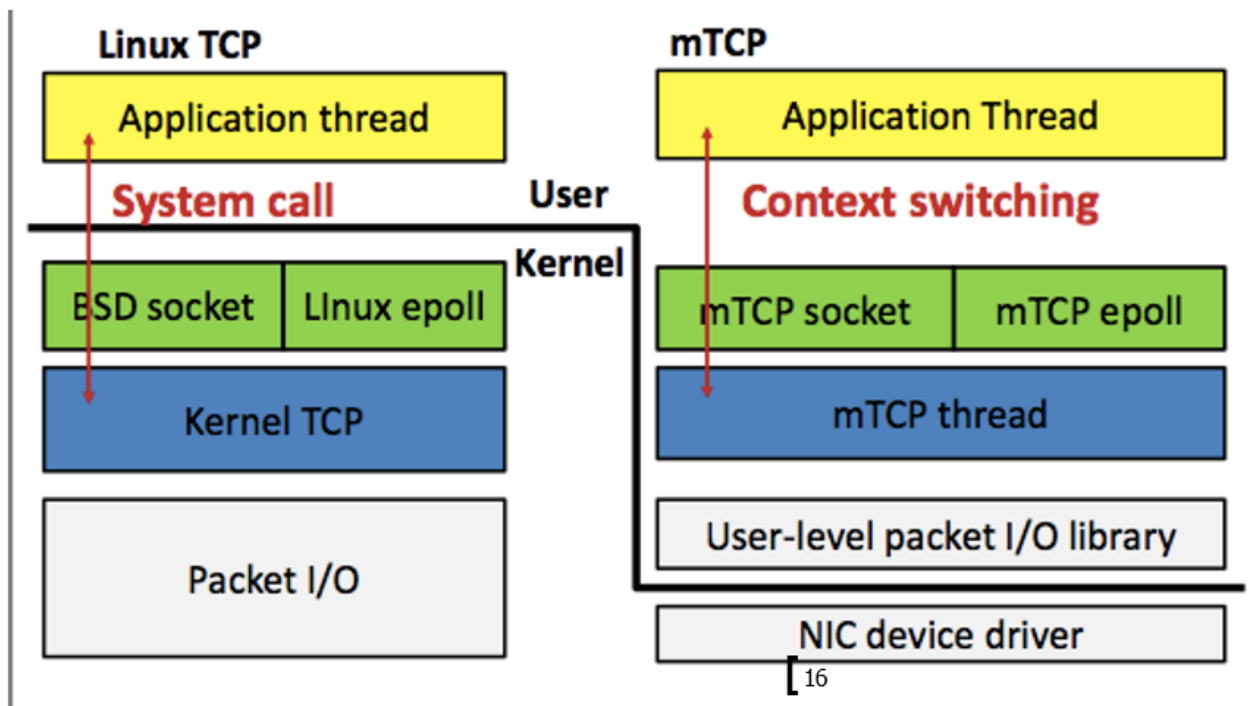
Solution:

- Batching in packet I/O, TCP processing, user applications (reduce system call overhead)
- Connection locality on multicore systems - handling same connection on same core, avoid cache pollution (solve connection locality)
- No descriptor sharing between mTCP thread

<PROJECT DATA SPACE>

mTCP: A Highly Scalable User-level TCP Stack for Multicore Systems <https://github.com/eunyoung14/mtcp>

mTCP Design



clone of mTCP in ES codeshare <http://git.es.f5net.com/index.cgi/codeshare/tree/vli/mtcp>

clone addition:

- Change apachebench configure script to compile with dpdk support
- Ported SSL BIO onto mTCP to enable apachebench to perform SSL test
- Add SSL clienthello stress test based on epwget and ssl-dos

- Add command line option in epwget and apachebench to enable source address pool to congest servers
- Increase mTCP SYN BACKLOG to increase concurrent connection
- Changed DPDK .config to compile DPDK as combined shared library
- Tuned send/receive buffer size in epwget.conf to achieve ~7 million concurrent connection on Dell Poweredge R710 II 72G MEM, 16 core, Intel NIC 82599ES

mTCP installation

<https://github.com/eunyoung14/mtcp> has detail installation

- DPDK VERSION -

1. Set up Intel's DPDK driver. Please use our version of DPDK.
We have only changed the lib/igb_uio/ submodule. The best method to compile DPDK package is to use DPDK's tools/setup.sh script. Please compile your package based on your own hardware configuration. We tested the mTCP stack on Intel Xeon E5-2690 (x86_64) machine with Intel 82599 Ethernet adapters (10G). We used the following steps in the setup.sh script for our setup:
 - Press [10] to compile the package
 - Press [13] to install the driver
 - Press [17] to setup 1024 2MB hugepages
 - Press [19] to register the Ethernet ports
 - Press [31] to quit the tool
 - check that DPDK package creates a new directory of compiled libraries. For x86_64 machines, the new subdirectory should be *dpdk-2.1.0/x86_64-native-linuxapp-gcc*
 - only those devices will work with DPDK drivers that are listed on this page: <http://dpdk.org/doc/nics>. Please make sure that your NIC is compatible before moving on to the next step.
2. Next bring the dpdk-registered interfaces up. Please use the setup_iface_single_process.sh script file present in dpdk-2.1.0/tools/ directory for this purpose. Please change lines 49-51 to change the IP address. Under default settings, run the script as:
./setup_iface_single_process.sh 3

This sets the IP address of your interfaces as 10.0.x.3.

3. Create soft links for include/ and lib/ directories inside empty dpdk/ directory:
cd dpdk/
ln -s <path_to_dpdk_2_1_0_directory>/x86_64-native-linuxapp-gcc/lib lib
ln -s <path_to_dpdk_2_1_0_directory>/x86_64-native-linuxapp-gcc/include include
4. Setup mtcp library:
./configure --with-dpdk-lib=\$<path_to_mtcp_release_v3>/dpdk
And not dpdk-2.1.0!
e.g. ./configure --with-dpdk-lib=`echo \$PWD`/dpdk
cd mtcp/src
make
 - check libmtcp.a in mtcp/lib
 - check header files in mtcp/include
5. make in util/:
make
6. make in apps/example:
make
 - check example binary files
7. Check the configurations
 - epserver.conf for server-side configuration
 - epwget.conf for client-side configuration
 - you may write your own configuration file for your application
 - please see README.config for more details
 - for the latest version, dyanmic ARP learning is *DISABLED*
8. Run the applications!

mTCP App configuration

mtcp configuration file

The underlying I/O module you want to use. Please

```
# enable only one out of the two.

io = dpdk

num_cores = 8

num_ip = 64

# Number of memory channels per processor socket (dpdk-only)

num_mem_ch = 4

#----- DPDK ports -----#

#port = dpdk0 dpdk1

port = dpdk0

#port = dpdk0:0

#port = dpdk0:1


# Enable multi-process support (under development)

#multiprocess = 0 master

#multiprocess = 1


# Receive buffer size of sockets

rcvbuf = 512

# Send buffer size of sockets

sndbuf = 512

# Maximum concurrency per core

max_concurrency = 1000000

# Maximum number of socket buffers per core

# Set this to small value if there are many idle connections

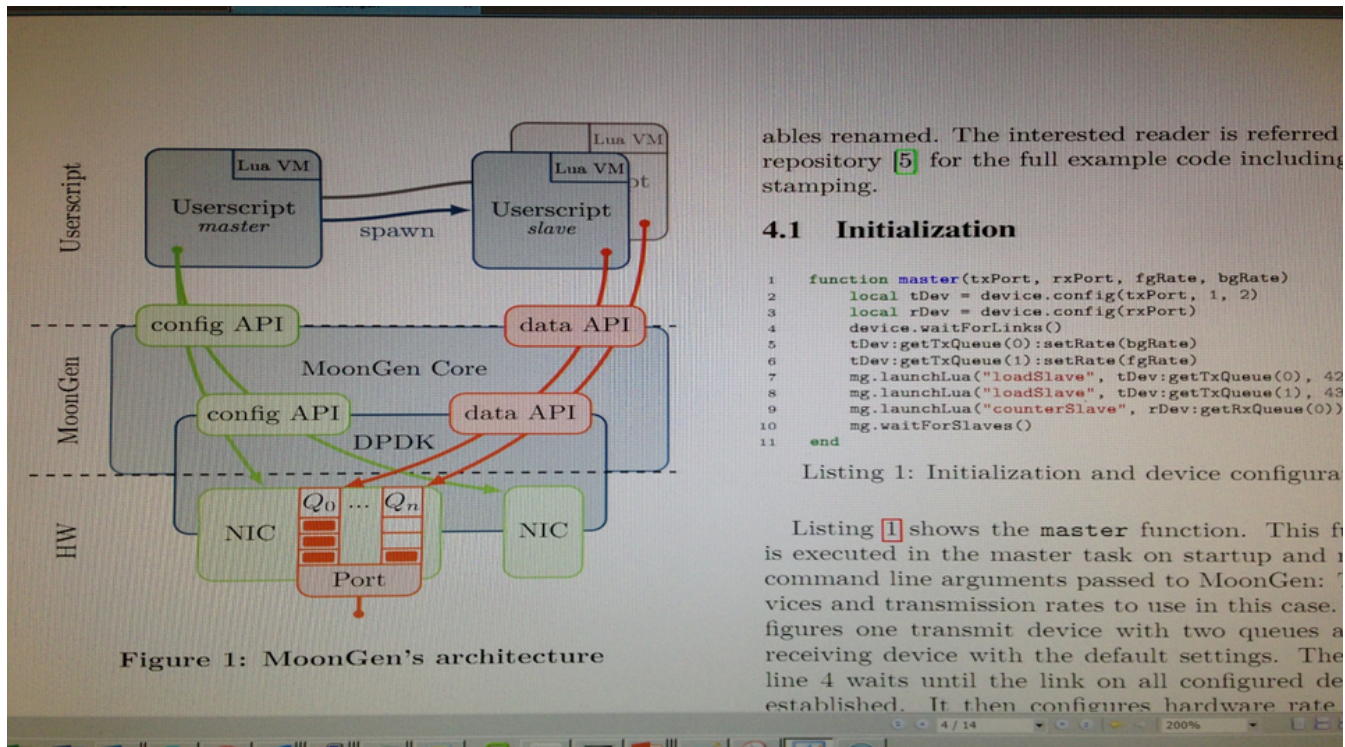
max_num_buffers = 1000000

# TCO timeout seconds

# (tcp_timeout = -1 can disable the timeout check)

tcp_timeout = 30
```

MoonGen: fully scriptable high-speed packet generator built on DPDK and LuaJIT. <https://github.com/emmericp/MoonGen>



ables renamed. The interested reader is referred repository [5] for the full example code including stamping.

4.1 Initialization

```

1  function master(txPort, rxPort, fgRate, bgRate)
2      local tDev = device.config(txPort, 1, 2)
3      local rDev = device.config(rxPort)
4      device.waitForLinks()
5      tDev:getTxQueue(0):setRate(bgRate)
6      tDev:getTxQueue(1):setRate(fgRate)
7      mg.launchLua("loadSlave", tDev:getTxQueue(0), 42
8      mg.launchLua("loadSlave", tDev:getTxQueue(1), 43
9      mg.launchLua("counterSlave", rDev:getRxQueue(0))
10     mg.waitForSlaves()
11 end

```

Listing 1: Initialization and device configura

Listing 1 shows the **master** function. This f is executed in the master task on startup and i command line arguments passed to MoonGen: r vices and transmission rates to use in this case. figures one transmit device with two queues a receiving device with the default settings. The line 4 waits until the link on all configured de established. It then configures hardware rate.

Userscript Master

```

1  local dpdk = require "dpdk"
2  local memory = require "memory"
3  local device = require "device"
4  local stats = require "stats"
5
6
7  function master(txPorts, minIp, numIps, rate)
8      if not txPorts then
9          printf("usage: txPort1[,txPort2[,...]] [minIP numIPs rate]")
10         return
11     end
12     txPorts = tostring(txPorts)
13     minIp = minIp or "10.0.0.1"
14     numIps = numIps or 100
15     rate = rate or 0
16     for currentTxPort in txPorts:gmatch("(%d+),?") do
17         currentTxPort = tonumber(currentTxPort)
18         local txDev = device.config({ port = currentTxPort })
19         txDev:wait()
20         txDev:getTxQueue(0):setRate(rate)
21         dpdk.launchLua("loadSlave", currentTxPort, 0, minIp, numIps)
22     end
23     dpdk.waitForSlaves()
24 end
25

```

Userscript slave


```

26 function loadSlave(port, queue, minA, numIPs)
27     --- parse and check ip addresses
28
29     local minIP, ipv4 = parseIPAddress(minA)
30     if minIP then
31         printf("INFO: Detected an %s address.", minIP and "IPv4" or "IPv6")
32     else
33         errorf("ERROR: Invalid minIP: %s", minA)
34     end
35
36     -- min TCP packet size for IPv6 is 74 bytes (+ CRC)
37     local packetLen = ipv4 and 60 or 74
38
39     --continue normally
40     local queue = device.get(port):getTxQueue(queue)
41     --local srcport = math.random(0, 2^16 - 1)
42     local mem = memory.createMemPool(function(buf)
43         buf:getTcpPacket(ipv4):fill(
44             ethSrc="a0:36:9f:a1:4d:6d", ethDst="52:54:00:2E:62:A2",
45             --- ip4Dst="10.9.1.2",
46             ip4Dst="10.9.3.6",
47             ip6Dst="fd06::1",
48             tcpDst="80",
49             tcpSyn=1,
50             tcpSeqNumber=1,
51             tcpWindow=10,
52             pktLength=packetLen )
53     end)
54
55     local lastPrint = dpdk.getTime()
56     local totalSent = 0
57     local lastTotal = 0
58     local lastSent = 0
59     local bufs = mem:bufArray(128)
60     local counter = 0
61     local c = 0
62
63     local txStats = stats:newDevTxCounter(queue, "plain")

```

```

64     while dpdk.running() do
65         -- faill packets and set their size
66         bufs:alloc(packetLen)
67         for i, buf in ipairs(bufs) do
68             local pkt = buf:getTcpPacket(ipv4)
69
70             --increment IP
71             if ipv4 then
72                 pkt.ip4.src:set(minIP)
73                 pkt.ip4.src:add(counter)
74                 pkt.tcp.src = math.random(0, 2^16 - 1)
75             else
76                 pkt.ip6.src:set(minIP)
77                 pkt.ip6.src:add(counter)
78             end
79             counter = incAndWrap(counter, numIPs)
80
81             -- dump first 3 packets
82             if c < 3 then
83                 buf:dump()
84                 c = c + 1
85             end
86         end
87         --offload checksums to NIC
88         bufs:offloadTcpChecksums(ipv4)
89
90         totalSent = totalSent + queue:send(bufs)
91         txStats:update()
92     end
93     txStats:finalize()
94 end

```

clone of MoonGen in ES codeshare <http://git.es.f5net.com/index.cgi/codeshare/tree/vli/MoonGen>

- improved tcp syn flooding with random src ip and src port
- added DNS flooding script to test Victoria2 DNS DDOS Hardware protection
- added icmp echo flooding

MoonGen Installation:

1. Install the dependencies (see below)
2. git submodule update --init
3. ./build.sh
4. ./setup-hugetlbfs.sh
5. Run MoonGen from the build directory

How to Run MoonGen script:

command syntax: build/MoonGen examples/<scriptname> <dpdk port> <min src ip> <# of src ip> <rate>

#build/MoonGen examples/dns-flood-victoria.lua 0 10.0.0.1 16000000 10000

Hardware SPEC:

Dell Poweredge R710 72G MEM, 16 core, Intel NIC 82599
2.40GHz 20 cores 64G MEM

DUT: Victoria B2250 Intel(R) Xeon(R) CPU E5-2658 v2 @

Dell PowerEdge R210 II (used \$300) 8 core, 32G MEM Intel 1G NIC I350
(cpu64-rhel6) 4 cores 16G MEM

DUT: BIGIP KVM VE CPU: QEMU Virtual CPU version

Load Test Example

1, DNS flooding without HW acceleration

#build/MoonGen examples/dns-flood-victoria.lua 0 10.0.0.1 16000000 10000

Device: id=0] Sent 13710082176 packets, current rate 4.51 Mpps, 3246.01 MBit/s, 3967.35 MBit/s wire rate.

[Device: id=0] Sent 13714591360 packets, current rate 4.51 Mpps, 3246.53 MBit/s, 3967.98 MBit/s wire rate.

[Device: id=0] Sent 13719099520 packets, current rate 4.51 Mpps, 3245.79 MBit/s, 3967.08 MBit/s wire rate.

top - 12:07:02 up 1 day, 20:38, 1 user, load average: 5.22, 7.46, 9.27

Tasks: 777 total, 19 running, 758 sleeping, 0 stopped, 0 zombie

Cpu(s): 50.6%us, 40.2%sy, 0.0%ni, 9.2%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st

Mem: 66080376k total, 65722732k used, 357644k free, 108700k buffers

Swap: 5242872k total, 0k used, 5242872k free, 4048612k cached

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
17859	root	1	-19	57.9g	145m	123m	R	100.8	0.2	56:44.33	tmm.0 -T 10 --tmid <=====HIGH CPU
17741	root	1	-19	57.9g	145m	123m	R	100.5	0.2	58:08.51	tmm.0 -T 10 --tmid
17853	root	1	-19	57.9g	145m	123m	R	100.5	0.2	56:46.73	tmm.0 -T 10 --tmid
17854	root	1	-19	57.9g	145m	123m	R	100.5	0.2	56:46.97	tmm.0 -T 10 --tmid
17855	root	1	-19	57.9g	145m	123m	R	100.5	0.2	56:46.06	tmm.0 -T 10 --tmid
17856	root	1	-19	57.9g	145m	123m	R	100.5	0.2	56:37.67	tmm.0 -T 10 --tmid

```

17857 root    1 -19 57.9g 145m 123m R 100.5 0.2 56:45.54 tmm.0 -T 10 --tmid
17858 root    1 -19 57.9g 145m 123m R 100.5 0.2 56:45.70 tmm.0 -T 10 --tmid
17860 root    1 -19 57.9g 145m 123m R 100.5 0.2 56:45.65 tmm.0 -T 10 --tmid
17852 root    1 -19 57.9g 145m 123m R 100.2 0.2 56:50.91 tmm.0 -T 10 --tmid
20110 root    RT  0  0  0  0 S 80.6 0.0 0:27.55 [enforcer/11]
20111 root    RT  0  0  0  0 R 80.6 0.0 0:27.56 [enforcer/15]
20116 root    RT  0  0  0  0 R 80.6 0.0 0:27.50 [enforcer/13]
20108 root    RT  0  0  0  0 R 80.2 0.0 0:27.55 [enforcer/19]
20109 root    RT  0  0  0  0 R 80.2 0.0 0:27.57 [enforcer/17]
20112 root    RT  0  0  0  0 S 80.2 0.0 0:27.55 [enforcer/5]
20113 root    RT  0  0  0  0 R 80.2 0.0 0:27.52 [enforcer/1]

```

```

-----
Ltm::Virtual Server: vs_dns_10g
-----

```

Status

```

Availability : unknown
State       : enabled
Reason      : The children pool member(s) either don't have service check
CMP         : enabled
CMP Mode    : all-cpus
Destination : 10.3.3.249:53
PVA Acceleration : none

```

Traffic	ClientSide	Ephemeral	General
Bits In	11.5G	0	-
Bits Out	16.7G	0	-
Packets In	20.0M	0	-
Packets Out	20.0M	0	-
Current Connections	27.1M	0	-
Maximum Connections	27.1M	0	-
Total Connections	28.8M	0	-

2, DNS flooding with HW acceleration

```
#build/MoonGen examples/dns-flood-victoria.lua 0 10.0.0.1 16000000 10000
```

Device: id=0] Sent 13710082176 packets, current rate 4.51 Mpps, 3246.01 MBit/s, 3967.35 MBit/s wire rate.

[Device: id=0] Sent 13714591360 packets, current rate 4.51 Mpps, 3246.53 MBit/s, 3967.98 MBit/s wire rate.

[Device: id=0] Sent 13719099520 packets, current rate 4.51 Mpps, 3245.79 MBit/s, 3967.08 MBit/s wire rate.

<https://docs.f5net.com/display/PDDESIGN/DNS+DDoS+HW+Acceleration++Validation>

```
sys fpga firmware-config {
```

```
    type I7-intelligent-fpga
```

```

}

ltm profile dns /Common/dns_fpga {
    app-service none
    enable-hardware-query-validation yes
    enable-hardware-response-cache yes
}

ltm virtual /Common/vs_dns_10g {
    destination /Common/10.3.3.249:53
    ip-protocol udp
    mask 255.255.255.255
    profiles {
        /Common/dns_fpga { }
        /Common/udp_immediate { }
    }
    rules {
        /Common/dns_responder
    }
    source 0.0.0.0/0
    translate-address enabled
    translate-port enabled
}

```

top - 14:51:05 up 3:30, 1 user, load average: 0.12, 0.05, 0.01

Tasks: 771 total, 1 running, 770 sleeping, 0 stopped, 0 zombie

Cpu(s): 4.2%us, 0.5%sy, 0.0%ni, 95.2%id, 0.0%wa, 0.1%hi, 0.0%si, 0.0%st

Mem: 66080272k total, 63094488k used, 2985784k free, 61152k buffers

Swap: 5242876k total, 0k used, 5242876k free, 1352852k cached

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
6428	root	1	-19	58.4g	151m	122m	S	12.6	0.2	3:19.62	tmm.0 <=====TMM CPU usage drops significantly
6435	root	1	-19	58.4g	151m	122m	S	11.3	0.2	2:42.67	tmm.4
6432	root	1	-19	58.4g	151m	122m	S	10.9	0.2	2:44.57	tmm.1
6433	root	1	-19	58.4g	151m	122m	S	10.9	0.2	2:42.78	tmm.2
6434	root	1	-19	58.4g	151m	122m	S	10.9	0.2	2:40.69	tmm.3
6436	root	1	-19	58.4g	151m	122m	S	10.9	0.2	2:41.53	tmm.5
6437	root	1	-19	58.4g	151m	122m	S	10.9	0.2	2:42.68	tmm.6
6438	root	1	-19	58.4g	151m	122m	S	10.9	0.2	2:40.92	tmm.7
6439	root	1	-19	58.4g	151m	122m	S	10.9	0.2	2:41.87	tmm.8
6440	root	1	-19	58.4g	151m	122m	S	10.6	0.2	2:41.49	tmm.9
28351	root	-91	0	97592	81m	31m	S	2.0	0.1	7:00.29	bcmINTR
28589	root	20	0	97592	81m	31m	S	2.0	0.1	5:43.36	bcmCNTR.0

profile_dns_stat

name	vs_name	queries	drops	hw_malformed	hw_inspected	hw_cache_lookups
hw_cache_responses						

/Common/dns_fpga	/Common/vs_dns_10g	6981	0	0	29727032	29726590
						29720435

3, SYN flooding without hardware acceleration

#build/MoonGen examples/l3-tcp-syn-flood.lua 0 10.0.0.1 16000000 10000

[Device: id=0] Sent 7061632 packets, current rate 7.06 Mpps, 3615.47 MBit/s, 4745.31 MBit/s wire rate.

ltm profile fastl4 /Common/fl4_fpga {

app-service none

defaults-from /Common/fastL4

hardware-syn-cookie disabled

pva-offload-dynamic disabled

software-syn-cookie enabled

}

top - 10:53:51 up 42 min, 1 user, load average: 0.24, 0.23, 0.65

Tasks: 769 total, 10 running, 759 sleeping, 0 stopped, 0 zombie

Cpu(s): 35.4%us, 1.7%sy, 0.1%ni, 62.9%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st

Mem: 66080376k total, 62740700k used, 3339676k free, 45784k buffers

Swap: 5242872k total, 0k used, 5242872k free, 1199508k cached

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
19290	root	1	-19	57.9g	145m	123m	R	71.5	0.2	0:14.38	tmm.0 -T 10 --tmid <=====HIGH CPU
19291	root	1	-19	57.9g	145m	123m	R	70.2	0.2	0:14.53	tmm.0 -T 10 --tmid
19293	root	1	-19	57.9g	145m	123m	S	70.2	0.2	0:14.37	tmm.0 -T 10 --tmid
19267	root	1	-19	57.9g	145m	123m	R	69.8	0.2	0:30.32	tmm.0 -T 10 --tmid
19292	root	1	-19	57.9g	145m	123m	R	69.8	0.2	0:14.38	tmm.0 -T 10 --tmid
19298	root	1	-19	57.9g	145m	123m	R	69.8	0.2	0:14.72	tmm.0 -T 10 --tmid
19295	root	1	-19	57.9g	145m	123m	R	69.5	0.2	0:14.73	tmm.0 -T 10 --tmid
19296	root	1	-19	57.9g	145m	123m	R	69.5	0.2	0:14.03	tmm.0 -T 10 --tmid
19297	root	1	-19	57.9g	145m	123m	R	69.2	0.2	0:14.14	tmm.0 -T 10 --tmid
19294	root	1	-19	57.9g	145m	123m	R	65.2	0.2	0:13.31	tmm.0 -T 10 --tmid

4 SYN flooding with HW acceleration

#build/MoonGen examples/l3-tcp-syn-flood.lua 0 10.0.0.1 16000000 10000

[Device: id=0] Sent 7061632 packets, current rate 7.06 Mpps, 3615.47 MBit/s, 4745.31 MBit/s wire rate.

ltm profile fastl4 /Common/fl4_fpga {

app-service none

```

defaults-from /Common/fastL4
hardware-syn-cookie enabled
pva-offload-dynamic enabled
software-syn-cookie enabled
}

top - 10:50:08 up 38 min,  1 user,  load average: 0.06, 0.36, 0.81

Tasks: 769 total,  1 running, 768 sleeping,  0 stopped,  0 zombie

Cpu(s):  0.8%us,  0.2%sy,  0.0%ni, 98.5%id,  0.5%wa,  0.0%hi,  0.0%si,  0.0%st

Mem:  66080376k total, 62740552k used,  3339824k free,   45324k buffers

Swap: 5242872k total,    0k used, 5242872k free, 1199492k cached

  PID USER   PR  NI  VIRT  RES  SHR  S %CPU  %MEM   TIME+  COMMAND
19267 root    1  -19 57.9g 145m 123m S  3.6  0.2   0:11.87 tmm <=====CPU usage drops significantly with hardware
acceleration
19293 root    1  -19 57.9g 145m 123m S  1.3  0.2   0:01.72 tmm
19296 root    1  -19 57.9g 145m 123m S  1.3  0.2   0:01.36 tmm
19297 root    1  -19 57.9g 145m 123m S  1.3  0.2   0:01.37 tmm
19290 root    1  -19 57.9g 145m 123m S  1.0  0.2   0:01.38 tmm
19292 root    1  -19 57.9g 145m 123m S  1.0  0.2   0:01.73 tmm
19294 root    1  -19 57.9g 145m 123m S  1.0  0.2   0:01.35 tmm
19295 root    1  -19 57.9g 145m 123m S  1.0  0.2   0:02.12 tmm
19298 root    1  -19 57.9g 145m 123m S  1.0  0.2   0:02.11 tmm
19291 root    1  -19 57.9g 145m 123m S  0.7  0.2   0:01.73 tmm

```

5, 10M concurrent HTTP connection

#epwget 10.3.3.249/ 160000000 -N 16 -c 10000000

```

[CPU 0] dpdk0 flows: 625000, RX: 96382(pps) (err: 0), 0.10(Gbps), TX: 413888(pps), 0.64(Gbps)
[CPU 1] dpdk0 flows: 625000, RX: 101025(pps) (err: 0), 0.10(Gbps), TX: 398592(pps), 0.61(Gbps)
[CPU 2] dpdk0 flows: 625000, RX: 106882(pps) (err: 0), 0.11(Gbps), TX: 418432(pps), 0.64(Gbps)
[CPU 3] dpdk0 flows: 625000, RX: 101497(pps) (err: 0), 0.10(Gbps), TX: 405952(pps), 0.62(Gbps)
[CPU 4] dpdk0 flows: 625000, RX: 107375(pps) (err: 0), 0.11(Gbps), TX: 427008(pps), 0.66(Gbps)
[CPU 5] dpdk0 flows: 625000, RX: 96012(pps) (err: 0), 0.10(Gbps), TX: 404352(pps), 0.62(Gbps)
[CPU 6] dpdk0 flows: 625000, RX: 100834(pps) (err: 0), 0.10(Gbps), TX: 405504(pps), 0.62(Gbps)
[CPU 7] dpdk0 flows: 625000, RX: 102572(pps) (err: 0), 0.11(Gbps), TX: 401024(pps), 0.62(Gbps)
[CPU 8] dpdk0 flows: 635366, RX: 111319(pps) (err: 0), 0.12(Gbps), TX: 410880(pps), 0.63(Gbps)
[CPU 9] dpdk0 flows: 625000, RX: 102179(pps) (err: 0), 0.11(Gbps), TX: 391104(pps), 0.60(Gbps)
[CPU10] dpdk0 flows: 625000, RX: 98014(pps) (err: 0), 0.10(Gbps), TX: 408320(pps), 0.63(Gbps)
[CPU11] dpdk0 flows: 625000, RX: 102712(pps) (err: 0), 0.11(Gbps), TX: 398976(pps), 0.61(Gbps)
[CPU12] dpdk0 flows: 625000, RX: 105891(pps) (err: 0), 0.11(Gbps), TX: 415616(pps), 0.64(Gbps)

```

[CPU13] dpdk0 flows: 625000, RX: 97728(pps) (err: 0), 0.10(Gbps), TX: 390592(pps), 0.60(Gbps)
[CPU14] dpdk0 flows: 625001, RX: 100570(pps) (err: 0), 0.10(Gbps), TX: 407872(pps), 0.63(Gbps)
[CPU15] dpdk0 flows: 625000, RX: 103412(pps) (err: 0), 0.11(Gbps), TX: 391296(pps), 0.60(Gbps)
[ALL] dpdk0 flows: 10010366, RX: 1634404(pps) (err: 0), 1.69(Gbps), TX: 6489408(pps), 9.96(Gbps) <=====

top - 15:25:26 up 23:57, 1 user, load average: 0.16, 0.33, 0.43

Tasks: 778 total, 17 running, 761 sleeping, 0 stopped, 0 zombie

Cpu(s): 45.1%us, 30.6%sy, 0.0%ni, 24.3%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st

Mem: 66080376k total, 62855960k used, 3224416k free, 136316k buffers

Swap: 5242872k total, 0k used, 5242872k free, 1182216k cached

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
17283	root	1	-19	57.9g	145m	123m	R	94.1	0.2	1322:36	tmm.0 -T 10 --tmid <=====
17286	root	1	-19	57.9g	145m	123m	R	94.1	0.2	1322:37	tmm.0 -T 10 --tmid
17281	root	1	-19	57.9g	145m	123m	R	93.8	0.2	1322:39	tmm.0 -T 10 --tmid
17284	root	1	-19	57.9g	145m	123m	R	93.8	0.2	1322:37	tmm.0 -T 10 --tmid
17282	root	1	-19	57.9g	145m	123m	R	93.4	0.2	1322:37	tmm.0 -T 10 --tmid
17287	root	1	-19	57.9g	145m	123m	R	93.4	0.2	1322:37	tmm.0 -T 10 --tmid
17288	root	1	-19	57.9g	145m	123m	R	93.4	0.2	1322:36	tmm.0 -T 10 --tmid
17289	root	1	-19	57.9g	145m	123m	R	93.4	0.2	1322:36	tmm.0 -T 10 --tmid
17043	root	1	-19	57.9g	145m	123m	R	92.8	0.2	1325:48	tmm.0 -T 10 --tmid
17285	root	1	-19	57.9g	145m	123m	R	92.1	0.2	1322:37	tmm.0 -T 10 --tmid
31507	root	RT	0	0	0	0	R	32.0	0.0	0:00.97	[enforcer/19]
31508	root	RT	0	0	0	0	R	32.0	0.0	0:00.97	[enforcer/13]
31509	root	RT	0	0	0	0	R	32.0	0.0	0:00.97	[enforcer/15]
31510	root	RT	0	0	0	0	S	31.7	0.0	0:00.96	[enforcer/9]
31511	root	RT	0	0	0	0	S	31.7	0.0	0:00.96	[enforcer/7]
31512	root	RT	0	0	0	0	R	31.4	0.0	0:00.95	[enforcer/3]
31515	root	RT	0	0	0	0	S	16.8	0.0	0:00.51	[enforcer/1]

[root@localhost:/S1-green-P:Active:Standalone] config # tail -f /var/log/ltm

Nov 4 15:25:29 slot1/bigip1 warning tmm7[17043]: 011e0003:4: Aggressive mode sweeper: /Common/default-eviction-policy (70000000002d6) (global memory) 9864 Connections killed

Nov 4 15:25:29 slot1/bigip1 warning tmm7[17043]: 011e0002:4: sweeper_policy_bind_deactivation_update: Aggressive mode /Common/default-eviction-policy deactivated (70000000002d6) (global memory). (12793204/15051776 pages)

Nov 4 15:25:29 slot1/bigip1 warning tmm6[17043]: 011e0003:4: Aggressive mode sweeper: /Common/default-eviction-policy (60000000002d2) (global memory) 10122 Connections killed

Nov 4 15:25:29 slot1/bigip1 warning tmm6[17043]: 011e0002:4: sweeper_policy_bind_deactivation_update: Aggressive mode /Common/default-eviction-policy deactivated (60000000002d2) (global memory). (12792703/15051776 pages)

Nov 4 15:25:29 slot1/bigip1 warning tmm3[17043]: 011e0003:4: Aggressive mode sweeper: /Common/default-eviction-policy (30000000002de) (global memory) 10877 Connections killed

Nov 4 15:25:29 slot1/bigip1 warning tmm3[17043]: 011e0002:4: sweeper_policy_bind_deactivation_update: Aggressive mode /Common/default-eviction-policy deactivated (30000000002de) (global memory). (12787088/15051776 pages)

Nov 4 15:25:29 slot1/bigip1 warning tmm4[17043]: 011e0003:4: Aggressive mode sweeper: /Common/default-eviction-policy

(40000000002c2) (global memory) 10306 Connections killed

Nov 4 15:25:29 slot1/bigip1 warning tmm4[17043]: 011e0002:4: sweeper_policy_bind_deactivation_update: Aggressive mode /Common/default-eviction-policy deactivated (40000000002c2) (global memory). (12787088/15051776 pages)

Every 1.0s: tmsh show ltm virtual vs_http_10g Wed Nov 4 15:27:15 2015

Availability : unknown

State : enabled

Reason : The children pool member(s) either don't have service check

ing enabled, or service check results are not available yet

CMP : enabled

CMP Mode : all-cpus

Destination : 10.3.3.249:80

PVA Acceleration : none

Traffic	ClientSide	Ephemeral	General
Bits In	329.8G	0	-
Bits Out	90.4G	0	-
Packets In	287.6M	0	-
Packets Out	150.2M	0	-
Current Connections	6.1M	0	- <=====
Maximum Connections	6.7M	0	-
Total Connections	39.8M	0	-

mTCP perf top output ~70% cycles in Userspace

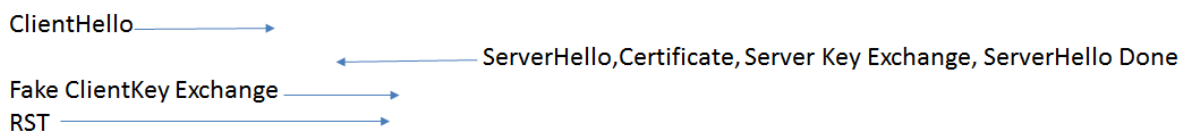
Samples: 1M of event 'cycles', Event count (approx.): 441906428558

8.25%	epwget	[.] SendTCPPacket
7.93%	[kernel]	[k] _raw_spin_lock
7.16%	epwget	[.] GetRSSCPUCore
7.15%	epwget	[.] IPOutput
4.26%	libc-2.19.so	[.] memset
4.10%	epwget	[.] ixgbe_xmit_pkts
3.62%	[kernel]	[k] clear_page_c
3.26%	epwget	[.] WriteTCPControlList
3.24%	[vdso]	[.] 0x00000000000000cf9
2.95%	epwget	[.] AddtoControlList
2.70%	epwget	[.] MTCPRunThread
2.66%	epwget	[.] HandleRTO
2.51%	epwget	[.] CheckRtmTimeout
2.10%	libpthread-2.19.so	[.] pthread_mutex_unlock
1.83%	epwget	[.] dpdk_send_pkts
1.68%	epwget	[.] HTInsert

1.65%	epwget	[.] CreateTCPStream
1.42%	epwget	[.] MPAllocateChunk
1.29%	epwget	[.] TCPCalcChecksum
1.24%	epwget	[.] dpdk_rcv_pkts
1.20%	epwget	[.] mtcp_getsockopt
1.12%	epwget	[.] rx_rcv_pkts

6, SSL DDOS test using mTCP

SSL DDOS using mTCP



```
#./apps/example/brute-shake 10.3.3.249 16000000 -N 16 -c 1600
```

```
[CPU 0] dpdk0 flows: 156, RX: 5733(pps) (err: 0), 0.04(Gbps), TX: 17208(pps), 0.03(Gbps)
```

```
.....
[CPU15] dpdk0 flows: 200, RX: 5674(pps) (err: 0), 0.04(Gbps), TX: 17034(pps), 0.03(Gbps)
```

```
[ ALL ] dpdk0 flows: 2795, RX: 60529(pps) (err: 0), 0.37(Gbps), TX: 178888(pps), 0.29(Gbps)
```

```
top - 09:10:21 up 22:58, 1 user, load average: 10.45, 4.43, 1.67
```

```
Tasks: 782 total, 19 running, 763 sleeping, 0 stopped, 0 zombie
```

```
Cpu(s): 50.6%us, 40.1%sy, 0.1%ni, 9.1%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
```

```
Mem: 66080376k total, 62923192k used, 3157184k free, 138624k buffers
```

```
Swap: 5242872k total, 0k used, 5242872k free, 1259132k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
21480	root	1	-19	57.9g	145m	123m	R	100.0	0.2	81:24.41	tmm <=====
21503	root	1	-19	57.9g	145m	123m	R	100.0	0.2	48:05.30	tmm
21504	root	1	-19	57.9g	145m	123m	R	100.0	0.2	47:23.12	tmm
21505	root	1	-19	57.9g	145m	123m	R	100.0	0.2	47:06.70	tmm
21506	root	1	-19	57.9g	145m	123m	R	100.0	0.2	46:55.21	tmm
21507	root	1	-19	57.9g	145m	123m	R	100.0	0.2	46:12.27	tmm
21508	root	1	-19	57.9g	145m	123m	R	100.0	0.2	46:56.27	tmm
21509	root	1	-19	57.9g	145m	123m	R	100.0	0.2	47:01.32	tmm
21510	root	1	-19	57.9g	145m	123m	R	100.0	0.2	46:48.54	tmm
21511	root	1	-19	57.9g	145m	123m	R	100.0	0.2	47:06.64	tmm
1675	root	RT	0	0	0	0	R	80.6	0.0	2:07.06	enforcer/9
1669	root	RT	0	0	0	0	R	80.2	0.0	2:07.03	enforcer/1
1670	root	RT	0	0	0	0	R	80.2	0.0	2:07.03	enforcer/15

```
1673 root    RT  0   0   0   0 R 80.2  0.0  2:07.03 enforcer/3
1677 root    RT  0   0   0   0 R 80.2  0.0  2:07.02 enforcer/13
1671 root    RT  0   0   0   0 S 79.9  0.0  2:07.04 enforcer/19
1672 root    RT  0   0   0   0 R 79.9  0.0  2:07.02 enforcer/5
```

profile_clientssl_stat

name	vs_name	cur_conns	max_conns	encrypted_bytes_in	encrypted_bytes_out	sess_cache_cur_entries

/Common/clientssl	/Common/vs_https	16235	25533	31164613471	66885162672	2621439

sess_cache_lookups	sess_cache_overflows	fatal_alerts	fully_hw_accelerated_conns	partially_hw_accelerated_conns		

47843605	42545150	55576	0	0		

SR 1-3391074511/#1-1K40BX7 example

```
root@pktgen-template:/home/admin/mtcp# ./apps/example/brute-shake 10.1.72.69 16000000 -N 8 -c 12000
```

7, ApacheBench(ab) mTCP port https test to Victoria

#ab -n 16000 -N 16 -c 8000 -L 64 https://10.3.3.249/

Loading mtcp configuration from : /etc/mtcp/config/mtcp.conf

Loading interface setting

EAL: Detected lcore 0 as core 0 on socket 0

.....

Checking link statusdone

Port 0 Link Up - speed 10000 Mbps - full-duplex

Benchmarking 10.3.3.249 (be patient)

CPU6 connecting to port 443

CPU7 connecting to port 443

CPU8 connecting to port 443

CPU9 connecting to port 443

CPU10 connecting to port 443

CPU5 connecting to port 443

CPU11 connecting to port 443

CPU12 connecting to port 443

CPU13 connecting to port 443

CPU14 connecting to port 443

CPU15 connecting to port 443

CPU4 connecting to port 443

CPU2 connecting to port 443
CPU3 connecting to port 443
CPU1 connecting to port 443
CPU0 connecting to port 443

[ALL] dpdk0 flows: 5016, RX: 9651(pps) (err: 0), 0.04(Gbps), TX: 14784(pps), 0.02(Gbps)

Ltm::Virtual Server: vs_https

CMP : enabled
CMP Mode : all-cpus
Destination : 10.3.3.249:443
PVA Acceleration : none

Traffic	ClientSide	Ephemeral	General
Bits In	49.2G	0	-
Bits Out	71.0G	0	-
Packets In	47.1M	0	-
Packets Out	30.4M	0	-
Current Connections	6.3K	0	-
Maximum Connections	146.0K	0	-
Total Connections	4.3M	0	-

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
12864	root	1	-19	57.9g	145m	123m	S	5.0	0.2	53:09.44	tm
13087	root	1	-19	57.9g	145m	123m	S	3.0	0.2	14:01.00	tm
13088	root	1	-19	57.9g	145m	123m	S	3.0	0.2	13:32.00	tm
13091	root	1	-19	57.9g	145m	123m	S	3.0	0.2	13:25.59	tm
13093	root	1	-19	57.9g	145m	123m	S	3.0	0.2	13:34.57	tm
13094	root	1	-19	57.9g	145m	123m	S	3.0	0.2	13:46.66	tm
13086	root	1	-19	57.9g	145m	123m	S	2.6	0.2	14:09.38	tm
13089	root	1	-19	57.9g	145m	123m	S	2.6	0.2	13:42.05	tm
13090	root	1	-19	57.9g	145m	123m	S	2.6	0.2	13:47.88	tm
13092	root	1	-19	57.9g	145m	123m	S	2.3	0.2	13:40.11	tm

8, ApacheBench(ab) mTCP port https test to BIGIP VE (KVM)

#ab -n 1000000 -c 8000 -N 8 -L 64 https://10.9.3.6/

Checking link status.....done
Port 0 Link Up - speed 1000 Mbps - full-duplex
Benchmarking 10.9.3.6 (be patient)

CPU6 connecting to port 443

CPU7 connecting to port 443

CPU5 connecting to port 443

CPU4 connecting to port 443

CPU3 connecting to port 443

CPU2 connecting to port 443

CPU1 connecting to port 443

CPU0 connecting to port 443

[ALL] dpdk0 flows: 8000, RX: 13443(pps) (err: 0), 0.01(Gbps), TX: 13953(pps), 0.01(Gbps)

top - 13:12:22 up 4 min, 1 user, load average: 3.34, 2.01, 0.82

Tasks: 395 total, 4 running, 391 sleeping, 0 stopped, 0 zombie

Cpu(s): 13.2%us, 6.5%sy, 0.0%ni, 64.5%id, 15.6%wa, 0.0%hi, 0.1%si, 0.0%st

Mem: 14403128k total, 14060912k used, 342216k free, 22400k buffers

Swap: 1048568k total, 0k used, 1048568k free, 863780k cached

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	P	COMMAND
13954	root	RT	0	12.0g	124m	104m	R	92.4	0.9	0:27.17	0	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088
14125	root	RT	0	12.0g	124m	104m	R	92.0	0.9	0:13.28	1	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088
14126	root	RT	0	12.0g	124m	104m	S	92.0	0.9	0:12.36	2	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088
14127	root	RT	0	12.0g	124m	104m	S	92.0	0.9	0:13.15	3	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088

Ltm::Virtual Server: vs_https

Status

Traffic	ClientSide	Ephemeral	General
Bits In	428.8M	0	-
Bits Out	786.2M	0	-
Packets In	505.9K	0	-
Packets Out	423.1K	0	-
Current Connections	9.4K	0	-
Maximum Connections	12.9K	0	-

9 Generate TCP/HTTP connection with random src MAC

<http://sourceforge.net/p/curl-loader/mailman/message/33614941/>

```

root@pktgen:/usr/src/mtcp# git diff mtcp/src/eth_out.c
diff --git a/mtcp/src/eth_out.c b/mtcp/src/eth_out.c
index 7fd1097..b064e9c 100644
--- a/mtcp/src/eth_out.c
+++ b/mtcp/src/eth_out.c
@@ -181,6 +181,44 @@ FlushSendChunkBuf(mtcp_manager_t mtcp, int nif)
 }
#endif /* E_PSIO */
/*-----*/
+/**
+ * Get a pseudo-random value.
+ *
+ * This function generates pseudo-random numbers using the linear
+ * congruential algorithm and 48-bit integer arithmetic, called twice
+ * to generate a 64-bit value.
+ *
+ * @return
+ *   A pseudo-random value between 0 and (1<<64)-1.
+ */
+static inline uint64_t
+rte_rand(void)
+{
+    uint64_t val;
+    val = lrand48();
+    val <<= 32;
+    val += lrand48();
+    return val;
+}
+
+static void
+generate_random_mac_addr(struct ethhdr *mac_addr)
+{
+    uint64_t random;
+
+    /* Set Organizationally Unique Identifier (OUI) prefix. */
+    mac_addr->h_source[0] = 0xa0;
+    mac_addr->h_source[1] = 0x36;
+    mac_addr->h_source[2] = 0x9f;
+    /* Force indication of locally assigned MAC address. */
+    mac_addr->h_source[0] |= ETHER_LOCAL_ADMIN_ADDR;
+    /* Generate the last 3 bytes of the MAC address with a random number. */
+    random = rte_rand();
+    memcpy(&mac_addr->h_source[3], &random, 3);
+}
+
uint8_t *
EthernetOutput(struct mtcp_manager *mtcp, uint16_t h_proto,
               int nif, unsigned char* dst_haddr, uint16_t iplen)
@@ -213,8 +251,9 @@ EthernetOutput(struct mtcp_manager *mtcp, uint16_t h_proto,
#endif

    ethh = (struct ethhdr *)buf;
+    generate_random_mac_addr(ethh);
    for (i = 0; i < ETH_ALEN; i++) {
        ethh->h_source[i] = CONFIG.eths[nif].haddr[i];
        //ethh->h_source[i] = CONFIG.eths[nif].haddr[i];
        ethh->h_dest[i] = dst_haddr[i];
    }
    ethh->h_proto = htons(h_proto);

```

Filter: tcp.flags.syn eq 1 and tcp.flags.ack ne 1											Expression... Clear Apply Save		
Index	Chain	Match	Jump	Source	Destination	Port	Port	Protocol	Length	Info			
8, 240858, 4093	IntelCor_4074:139	RealtekU_2e:62:a2	10.9.3.9	10.9.3.6	1026	80		TCP	158	IN s1/tm2 : 1026-80 [SYN] Seq=1079055891 Win=14600 Len=0 MSS=1460 TSval=147454791 TSecr=0			
0, 000018, 4093	IntelCor_21:2d:1f4	RealtekU_2e:62:a2	10.9.3.9	10.9.3.6	1030	80		TCP	158	IN s1/tm2 : 1030-80 [SYN] Seq=1395331858 Win=14600 Len=0 MSS=1460 TSval=147454791 TSecr=0			
0, 001442, 4093	IntelCor_1c:a9:c1	RealtekU_2e:62:a2	10.9.3.9	10.9.3.6	1034	80		TCP	158	IN s1/tm2 : 1034-80 [SYN] Seq=1891673374 Win=14600 Len=0 MSS=1460 TSval=147454799 TSecr=0			
0, 000016, 4093	IntelCor_99:2d:193	RealtekU_2e:62:a2	10.9.3.9	10.9.3.6	1038	80		TCP	158	IN s1/tm2 : 1038-80 [SYN] Seq=2139790996 Win=14600 Len=0 MSS=1460 TSval=147454799 TSecr=0			

10 mTCP hack to generate TCP 3WHS final ACK with 1 byte payload

SR 1-2235587071/ IXIA were sending 1byte payload in TCP 3WHS final ACK to cause BIGIP memory sweeper kicking in, unable to reach advertised platform tcp connection spec numbers, mTCP patch below to simulate that

```
oot@pktgen:/home/dpdk/mtcp# git diff mtcp/src/tcp_out.c
diff --git a/mtcp/src/tcp_out.c b/mtcp/src/tcp_out.c
index e346958..edebe19 100644
--- a/mtcp/src/tcp_out.c
+++ b/mtcp/src/tcp_out.c
@@ -432,6 +432,8 @@ SendControlPacket(mtcp_manager_t mtcp, tcp_stream *cur_stream, uint32_t cur_ts)
{
    struct tcp_send_vars *sndvar = cur_stream->sndvar;
    int ret = 0;
    char payload = 'y';
    uint16_t payloadlen = 1;

    if (cur_stream->state == TCP_ST_SYN_SENT) {
        /* Send SYN here */
@@ -445,7 +447,8 @@ SendControlPacket(mtcp_manager_t mtcp, tcp_stream *cur_stream, uint32_t cur_ts)

    } else if (cur_stream->state == TCP_ST_ESTABLISHED) {
        /* Send ACK here */
        ret = SendTCPacket(mtcp, cur_stream, cur_ts, TCP_FLAG_ACK, NULL, 0);
//        ret = SendTCPacket(mtcp, cur_stream, cur_ts, TCP_FLAG_ACK, NULL, 0);
        ret = SendTCPacket(mtcp, cur_stream, cur_ts, TCP_FLAG_ACK, (uint8_t *)&payload, payloadlen);
}
```

Jul 18 14:08:47 slot1/cluster1 warning tmm3[22490]: 011e0002:4: sweeper_policy_bind_deactivation_update: Aggressive mode /Common/default-eviction-policy deactivated (3000000000001) (global memory). (1524040/1793024 pages)

Jul 18 14:08:47 slot1/cluster1 warning tmm[22490]: 011e0003:4: Aggressive mode sweeper: /Common/default-eviction-policy (1) (global memory) 20244 Connections killed

Jul 18 14:08:47 slot1/cluster1 warning tmm[22490]: 011e0002:4: sweeper_policy_bind_deactivation_update: Aggressive mode /Common/default-eviction-policy deactivated (1) (global memory). (1524010/1793024 pages)

filter	3.2G	3.6G	1
umem	1.2G	10.2G	1
xdata	4.0G	4.6G	2048

11 ICMP ping flooding to BIGIP VE

build/MoonGen examples/icmp-flood.lua 0 10.0.0.1 16000000 10000

top - 12:10:54 up 1:55, 1 user, load average: 0.24, 0.06, 0.02

Tasks: 381 total, 2 running, 379 sleeping, 0 stopped, 0 zombie

Cpu0 : 16.2%us, 11.3%sy, 0.0%ni, 13.1%id, 0.0%wa, 0.3%hi, 59.1%si, 0.0%st

Cpu1 : 2.1%us, 2.4%sy, 0.0%ni, 94.6%id, 0.0%wa, 0.0%hi, 0.3%si, 0.6%st

Cpu2 : 3.5%us, 3.2%sy, 0.0%ni, 90.6%id, 0.0%wa, 0.0%hi, 1.8%si, 0.9%st

Cpu3 : 1.2%us, 1.8%sy, 0.3%ni, 96.0%id, 0.0%wa, 0.0%hi, 0.3%si, 0.3%st

Mem: 14403128k total, 14267112k used, 136016k free, 22252k buffers

Swap: 1048568k total, 1224k used, 1047344k free, 571908k cached

PID USER PR NI VIRT RES SHR S %CPU %MEM TIME+ P COMMAND

2889 root	RT	0	12.0g	124m	104m	S	154.7	0.9	3:58.44	0	tmm.0	-T 4	--tmid 0	--npus 4	--platform Z100	-m -s 12088
3054 root	RT	0	12.0g	124m	104m	S	8.8	0.9	2:48.53	3	tmm.0	-T 4	--tmid 0	--npus 4	--platform Z100	-m -s 12088
3053 root	RT	0	12.0g	124m	104m	S	8.5	0.9	2:12.29	2	tmm.0	-T 4	--tmid 0	--npus 4	--platform Z100	-m -s 12088
3052 root	RT	0	12.0g	124m	104m	R	7.6	0.9	2:06.63	1	tmm.0	-T 4	--tmid 0	--npus 4	--platform Z100	-m -s 12088

Technical tips for load generation

1 mTCP thread pre-allocate memory pools for TCP send and receive buffer from configured maximum number of buffers.

when the load generator has limited memory, it is recommended to reduce the size of TCP send and receive buffer and number of buffers in application configuration file.

also configure the BIGIP DUT to respond small packet size (< 64 bytes) because large response payload size would trigger mTCP payload merge length error

for example: setting TCP receive and send buffer in epwget.conf

```
# Receive buffer size of sockets
```

```
rcvbuf = 1024
```

```
# Send buffer size of sockets
```

```
sndbuf = 1024
```

2 mTCP default receive window is 64 bytes when window scale not negotiated, TMOS 11.4.0 and later disabled window scale unless increase send/receive buffer > 65535

This result in BIGIP send 64 bytes to mTCP in each segment while doing load testing, cause BIGIP buffers data in longer period and connection stays active longer period

see [SR 1-1526064281/](#). Linux client would not notice this difference when window scale not negotiated because Linux client advertise receive window > MSS

3 When BIGIP syncookie activated, there are tons of TCP retransmission between mTCP and BIGIP, filed an issue with mTCP <https://github.com/eunyoung14/mtcp/issues/23>

also BIGIP bug is filed in [Bug 554761](#)

4 Victoria2 DNS DDOS Hardware validation skips # of question check, [Bug 558251](#) DNS query with multiple question RR marked as malformed in SW and not in HW, without HW validation of # of DNS question count, TMM CPU usage tops at 100% under the DNS flooding test

5 BIGIP v12.0.0 SW syncookies and window scaling will cause 3WHS to fail on L7 VIP, this would also affect tcp/http load test [Bug 555020](#)

6 SSL performance test with modified multithread apachebench with SSL support [SR 1-1602453049/](#) (use AES256-SHA to get better performance number, not TLS_DHE_RSA_WITH_*)

7 Bug 486688 SYN cookie learning for network and/or wildcard port virtual servers : test code can be found here <http://ickernel.blogspot.com/2016/01/syn-flood-network-virtual-server-in.html>

8 ID 621988 SR 1-2402757471 same src ipport udp packet drops on udp virtual with mirror

9 DPDK + mTCP on VMware ESXi VM and Linux perf to resolve [SR 1-2653101191](#)

Development:

1. Ported mTCP/DPDK to VMware ESXi VM ([patch](https://docs.f5net.com/download/attachments/290265627/ESXi-VM-mTCP.txt?api=v2) <https://docs.f5net.com/download/attachments/290265627/ESXi-VM-mTCP.txt?api=v2>)

mTCP/DPDK in VM (vli_ubuntu_1604_DPDK):

(Note: Ubuntu 1604 changes mTCP DPDK netdev name dpdk0 once igb_uio bounded, to disable it, add GRUB_CMDLINE_LINUX="net.ifnames=0" and run update-grub)



Figure 3. Standard para-virtual device connectivity

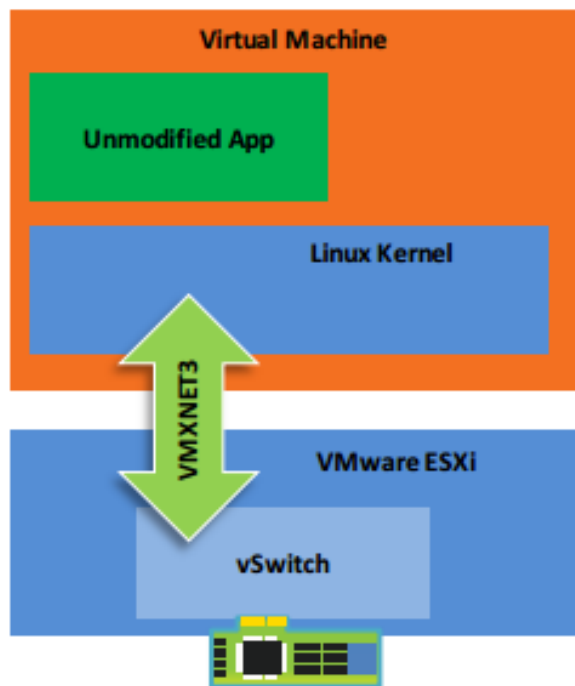
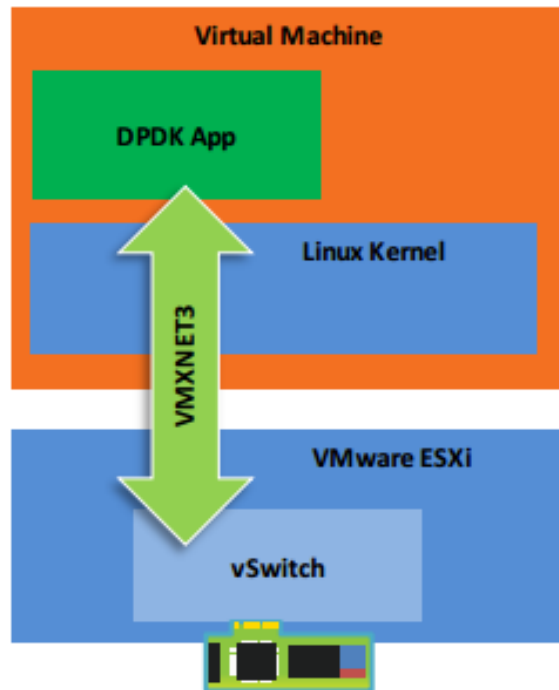


Figure 4. Intel DPDK para-virtual device connectivity



```
root@ubuntu-dpdk:/home/dpdk/mtcp# ./apps/example/epwget 10.169.72.88 1600000000 -N 8 -c 1600000
```

```
-----  
Loading mtcp configuration from : /etc/mtcp/config/epwget.conf  
Loading interface setting
```

```
EAL: PCI device 0000:03:00.0 on NUMA socket -1
```

```
EAL: probe driver: 15ad:7b0 rte_vmxnet3_pmd
```

```
EAL: Not managed by a supported kernel driver, skipped
```

```
EAL: PCI device 0000:0b:00.0 on NUMA socket -1
```

```
EAL: probe driver: 15ad:7b0 rte_vmxnet3_pmd
```

```
Interface name: dpdk0
```

```
Configurations:
```

```
Number of CPU cores available: 8
```

```
Number of CPU cores to use: 8
```

```
Number of source ip to use: 32
```

```
Maximum number of concurrency per core: 1000000
```

```
Maximum number of preallocated buffers per core: 1000000
```

```
Receive buffer size: 2048
```

```
Send buffer size: 2048
```

```
TCP timeout seconds: 30
```

```
TCP timewait seconds: 0
```

```
NICs to print statistics: dpdk0  
-----
```

Interfaces:

name: dpdk0, ifindex: 0, hwaddr: 00:50:56:86:43:CF, ipaddr: 10.0.0.1, netmask: 255.255.255.0

Number of NIC queues: 8

[CPU 0] dpdk0 flows: 200194, RX: 30694(pps) (err: 0), 0.03(Gbps), TX: 61504(pps), 0.09(Gbps)

[CPU 1] dpdk0 flows: 200114, RX: 30832(pps) (err: 0), 0.02(Gbps), TX: 61760(pps), 0.09(Gbps)

[CPU 2] dpdk0 flows: 200059, RX: 31250(pps) (err: 0), 0.03(Gbps), TX: 61696(pps), 0.09(Gbps)

[CPU 3] dpdk0 flows: 200030, RX: 31149(pps) (err: 0), 0.02(Gbps), TX: 61632(pps), 0.09(Gbps)

[CPU 4] dpdk0 flows: 200010, RX: 31682(pps) (err: 0), 0.03(Gbps), TX: 61632(pps), 0.09(Gbps)

[CPU 5] dpdk0 flows: 200005, RX: 31424(pps) (err: 0), 0.03(Gbps), TX: 61696(pps), 0.09(Gbps)

[CPU 6] dpdk0 flows: 200000, RX: 32209(pps) (err: 0), 0.03(Gbps), TX: 61824(pps), 0.09(Gbps)

[CPU 7] dpdk0 flows: 200001, RX: 45255(pps) (err: 0), 0.04(Gbps), TX: 61632(pps), 0.09(Gbps)

[ALL] dpdk0 flows: 1600413, RX: 264495(pps) (err: 0), 0.22(Gbps), TX: 493376(pps), 0.74(Gbps)

BIGIP 10200 -



2 Port MoonGen/DPDK to VMware ESXi VM (MoonGen DPDK is not synced with Intel DPDK, running MoonGen on VM can be challenging, Developed syn, dns flooding in DPDK Pktgen to replace MoonGen)

3 Intel Pktgen-DPDK in VMware ESXi VM (patches can found in attachment syn-dns-flood.txt)

a: syn flooding

some comments in SR 1-2068544601/#1-ZO4P5N

```
root@pktgen-template:/home/admin/pktgen-dpdk/dpdk/examples/pktgen-dpdk# ./app/app/x86_64-native-linuxapp-gcc/pktgen -c ff --  
-P -m [0:1-7].0
```

```
Pktgen> load synflood.txt
```

```
Pktgen>page range
```

```
Port #      Port-0  
dst.ip      : 10.2.72.97  
inc         : 0.0.0.1  
min         : 10.2.72.97  
max         : 10.2.72.200  
:
```

```

src.ip      :    10.1.72.154
inc         :    0.0.0.1
min         :    10.1.72.154
max         :    10.1.72.254
:
ip_proto    :    TCP
:
dst.port / inc :    80/ 0
min / max   :    0/ 0
:
src.port / inc :   1025/ 1
min / max   :   1025/65512
:
vlan.id / inc :    1/ 0
min / max   :    1/4095
:
pkt.size / inc :    64/ 0
min / max   :    64/1518
:
dst.mac     : 00:23:e9:6c:74:83
inc         : 00:00:00:00:00:00
min         : 00:00:00:00:00:00
max         : 00:00:00:00:00:00
:
src.mac     : 00:50:56:86:10:76
inc         : 00:00:00:00:00:00
min         : 00:00:00:00:00:00
max         : 00:00:00:00:00:00
:
gtpu.teid / inc :    0/ 0
min / max   :    0/ 0
-- Pktgen Ver: 2.9.17 (DPDK 16.04.0-rc2) Powered by Intel® DPDK -----

```

Pktgen> **start 0**

Pktgen> **stop 0**

```

root@pktgen-template:/home/admin/pktgen-dpdk/dpdk/examples/pktgen-dpdk# cat synflood.txt
#
# Pktgen - Ver: 2.9.17 (DPDK 16.04.0-rc2)
# Copyright (c) <2010-2016>, Intel Corporation. All rights reserved., Powered by Intel® DPDK

# Command line arguments: (DPDK args are defaults)
# ./app/app/x86_64-native-linuxapp-gcc/pktgen -c ff -n 3 -m 512 --proc-type primary -- -P -m [0:1-7].0

#####
# Pktgen Configuration script information:
# GUI socket is Not Enabled
# Flags 00040004
# Number of ports: 1
# Number ports per page: 4
# Number descriptors: RX 512 TX: 512
# Promiscuous mode is Enabled

#####
# Global configuration:
geometry 132x44
mac_from_arp disable

##### Port 0 #####
#
# Port: 0, Burst: 32, Rate:100%, Flags:c0000010, TX Count:Forever
# SeqCnt:0, Prime:1 VLAN ID:0001, Link: <UP-10000-FD>
#
# Set up the primary port information:
set 0 count 0
set 0 size 64
set 0 rate 100
set 0 burst 32

```

```
set 0 sport 1234
set 0 dport 5678
set 0 prime 1
type ipv4 0
proto tcp 0
set ip dst 0 10.2.72.97
#set ip dst 0 10.1.72.8
set ip src 0 10.1.72.154/24
#vpr 172.24.15.7
#set mac 0 00:23:E9:12:AA:01
#10200 172.24.40.6
#set mac 0 00:23:E9:E5:F2:C3
#4200 172.24.46.39
#set mac 0 00:23:E9:63:5B:83
#5200 172.24.19.15
set mac 0 00:23:E9:6C:74:83
vlanid 0 1
```

```
pattern 0 zero
user.pattern 0 0123456789abcdef
```

```
latency 0 disable
mpls 0 disable
mpls_entry 0 0
qinq 0 disable
qinqids 0 0 0
gre 0 disable
gre_eth 0 disable
gre_key 0 0
#
# Port flag values:
icmp.echo 0 disable
pcap 0 disable
range 0 enable
process 0 disable
capture 0 disable
rxtap 0 disable
txtap 0 disable
vlan 0 disable
```

```
#
# Range packet information:
src.mac start 0 00:50:56:86:10:76
src.mac min 0 00:00:00:00:00:00
src.mac max 0 00:00:00:00:00:00
src.mac inc 0 00:00:00:00:00:00
dst.mac start 0 00:23:E9:6C:74:83
#dst.mac start 0 00:23:E9:E5:F2:C3
#dst.mac start 0 00:23:E9:12:AA:01
#dst.mac start 0 00:23:E9:63:5B:83
#dst.mac start 0 00:50:56:86:84:90
dst.mac min 0 00:00:00:00:00:00
dst.mac max 0 00:00:00:00:00:00
dst.mac inc 0 00:00:00:00:00:00
```

```
src.ip start 0 10.1.72.154
src.ip min 0 10.1.72.154
src.ip max 0 10.1.72.254
src.ip inc 0 0.0.0.1
```

```
dst.ip start 0 10.2.72.97
dst.ip min 0 10.2.72.97
dst.ip max 0 10.2.72.200
dst.ip inc 0 0.0.0.1
```

```
#dst.ip start 0 10.1.72.8
#dst.ip min 0 10.1.72.8
#dst.ip max 0 10.1.72.8
#dst.ip inc 0 0.0.0.1
```

```
src.port start 0 1025
src.port min 0 1025
```

```
src.port max 0 65512
src.port inc 0 1
```

```
dst.port start 0 80
dst.port min 0 0
dst.port max 0 0
dst.port inc 0 0
```

```
vlan.id start 0 1
vlan.id min 0 1
vlan.id max 0 4095
vlan.id inc 0 0
```

```
pkt.size start 0 64
pkt.size min 0 64
pkt.size max 0 1518
pkt.size inc 0 0
```

```
#
# Set up the sequence data for the port.
set 0 seqCnt 0
```

```
##### Done #####
```

b: dns flooding

```
Pktgen> load dnsflood.txt
```

```
Pktgen> start 0
```

```
Pktgen> stop 0
```

```
root@pktgen-template:/home/admin/pktgen-dpdk/dpdk/examples/pktgen-dpdk# cat dnsflood.txt
```

```
#
# Pktgen - Ver: 2.9.17 (DPDK 16.04.0-rc2)
# Copyright (c) <2010-2016>, Intel Corporation. All rights reserved., Powered by Intel® DPDK

# Command line arguments: (DPDK args are defaults)
# ./app/app/x86_64-native-linuxapp-gcc/pktgen -c ff -n 3 -m 512 --proc-type primary -- -P -m [0:1-7].0
```

```
#####
```

```
# Pktgen Configuration script information:
# GUI socket is Not Enabled
# Flags 00040004
# Number of ports: 1
# Number ports per page: 4
# Number descriptors: RX 512 TX: 512
# Promiscuous mode is Enabled
```

```
#####
```

```
# Global configuration:
geometry 132x44
mac_from_arp disable
```

```
##### Port 0 #####
```

```
#
# Port: 0, Burst: 32, Rate:100%, Flags:c0000010, TX Count:Forever
# SeqCnt:0, Prime:1 VLAN ID:0001, Link: <UP-10000-FD>
#
# Set up the primary port information:
set 0 count 0
set 0 size 64
set 0 rate 100
set 0 burst 32
set 0 sport 1234
set 0 dport 5678
set 0 prime 1
type ipv4 0
proto udp 0
set ip dst 0 10.2.72.97
```

```
#set ip dst 0 10.1.72.8
set ip src 0 10.1.72.154/24
#vpr 172.24.15.7
#set mac 0 00:23:E9:12:AA:01
#10200 172.24.40.6
#set mac 0 00:23:E9:E5:F2:C3
#4200 172.24.46.39
#set mac 0 00:23:E9:63:5B:83
#5200 172.24.19.15
set mac 0 00:23:E9:6C:74:83
vlanid 0 1

#pattern 0 user
#user.pattern 0 00001111111111111111
```

```
latency 0 disable
mpls 0 disable
mpls_entry 0 0
qinq 0 disable
qinqids 0 0 0
gre 0 disable
gre_eth 0 disable
gre_key 0 0
#
# Port flag values:
icmp.echo 0 disable
pcap 0 disable
range 0 enable
range.proto 0 udp
process 0 disable
capture 0 disable
rxtap 0 disable
txtap 0 disable
vlan 0 disable
```

```
#
# Range packet information:
src.mac start 0 00:50:56:86:10:76
src.mac min 0 00:00:00:00:00:00
src.mac max 0 00:00:00:00:00:00
src.mac inc 0 00:00:00:00:00:00
dst.mac start 0 00:23:E9:6C:74:83
#dst.mac start 0 00:23:E9:E5:F2:C3
#dst.mac start 0 00:23:E9:12:AA:01
#dst.mac start 0 00:23:E9:63:5B:83
#dst.mac start 0 00:50:56:86:84:90
dst.mac min 0 00:00:00:00:00:00
dst.mac max 0 00:00:00:00:00:00
dst.mac inc 0 00:00:00:00:00:00
```

```
src.ip start 0 10.1.72.154
src.ip min 0 10.1.72.154
src.ip max 0 10.1.72.254
src.ip inc 0 0.0.0.1
```

```
dst.ip start 0 10.2.72.97
dst.ip min 0 10.2.72.97
dst.ip max 0 10.2.72.200
dst.ip inc 0 0.0.0.1
```

```
#dst.ip start 0 10.1.72.8
#dst.ip min 0 10.1.72.8
#dst.ip max 0 10.1.72.8
#dst.ip inc 0 0.0.0.1
```

```
src.port start 0 1025
src.port min 0 1025
src.port max 0 65512
src.port inc 0 1
```

```
dst.port start 0 53
dst.port min 0 0
```

```
dst.port max 0 0
dst.port inc 0 0
```

```
vlan.id start 0 1
vlan.id min 0 1
vlan.id max 0 4095
vlan.id inc 0 0
```

```
pkt.size start 0 74 #note packet size setting has to match the whole udp dns payload, can be improved in code to deal with variable length of
domain name query
```

```
pkt.size min 0 74
pkt.size max 0 1518
pkt.size inc 0 0
```

```
#
# Set up the sequence data for the port.
set 0 seqCnt 0
```

```
##### Done #####
```

Bug 486688 test validation

-----CLIENT VMware ESXi VM vli_ubuntu_1404_ab_template with DPDK pktgen

```
root@pktgen-template:/home/admin/pktgen-dpdk/dpdk/examples/pktgen-dpdk# ./app/app/x86_64-native-linuxapp-gcc/pktgen -c ff
-- -P -m "[0:0-7].0 "
```

Copyright (c) <2010-2016>, Intel Corporation. All rights reserved.
Pktgen created by: Keith Wiles -- >>> Powered by Intel® DPDK <<<

Lua 5.3.2 Copyright (C) 1994-2015 Lua.org, PUC-Rio
>>> Packet Burst 32, RX Desc 512, TX Desc 512, mbufs/port 4096, mbuf cache 512

```
=== port to lcore mapping table (# lcores 8) ===
lcore:  0  1  2  3  4  5  6  7
port  0: D: T 0: 1 0: 1 0: 1 0: 1 0: 1 0: 1 0: 1 = 1: 8
Total  : 1: 1 0: 1 0: 1 0: 1 0: 1 0: 1 0: 1 0: 1
Display and Timer on lcore 0, rx:tx counts per port/lcore
```

Configuring 1 ports, MBUF Size 1920, MBUF Cache Size 512

```
Lcore:
0, RX-TX
    RX( 1): ( 0: 0)
    TX( 1): ( 0: 0)
1, TX-Only
    TX( 1): ( 0: 1)
2, TX-Only
    TX( 1): ( 0: 2)
3, TX-Only
    TX( 1): ( 0: 3)
4, TX-Only
    TX( 1): ( 0: 4)
5, TX-Only
    TX( 1): ( 0: 5)
6, TX-Only
    TX( 1): ( 0: 6)
7, TX-Only
    TX( 1): ( 0: 7)
```

Port :
0, nb_lcores 8, private 0x8f0690, lcores: 0 1 2 3 4 5 6 7

```

** Dev Info (rte_vmxnet3_pmd:0) **
max_vfs      : 0 min_rx_bufsize :1646 max_rx_pktlen : 16384 max_rx_queues      : 16 max_tx_queues: 8
max_mac_addrs : 1 max_hash_mac_addrs: 0 max_vmdq_pools: 0
rx_offload_capa: 13 tx_offload_capa : 45 reta_size : 0 flow_type_rss_offloads:0000000000000514
vmdq_queue_base: 0 vmdq_queue_num : 0 vmdq_pool_base: 0
** RX Conf **
pthreash     : 0 hthresh      : 0 wthresh      : 0
Free Thresh  : 0 Drop Enable  : 0 Deferred Start : 0
** TX Conf **
pthreash     : 0 hthresh      : 0 wthresh      : 0
Free Thresh  : 0 RS Thresh    : 0 Deferred Start : 0 TXQ Flags:00000200

```

Initialize Port 0 -- TxQ 8, RxQ 1, Src MAC 00:50:56:86:10:76

Pktgen > **load synflood.txt**

Pktgen> **start 0**

Pktgen> **stop 0**

-----BIGIP DUT:

```

net vlan esnet-1101 {

```

```

    if-index 896

```

```

    interfaces {

```

```

        1.1 {

```

```

            tagged

```

```

        }

```

```

    }

```

```

    tag 1101

```

```

}

```

```

net vlan esnet-1102 {

```

```

    if-index 912

```

```

    interfaces {

```

```

        1.1 {

```

```

            tagged

```

```

        }

```

```

    }

```

```

    tag 1102

```

```

}

```

```

net self 10.2.72.12 {

```

```

    address 10.2.72.12/16

```

```

    allow-service all

```

```

    traffic-group traffic-group-local-only

```

```

    vlan esnet-1102

```

```

}

```

```

net self 10.1.72.12 {

```



```
address 10.1.72.12/16
allow-service all
traffic-group traffic-group-local-only
vlan esnet-1101
}
```

```
ltm virtual /Common/vs_http {
  destination /Common/10.2.0.0:80
  ip-forward
  ip-protocol tcp
  mask 255.255.0.0
  profiles {
    /Common/fl4_swsyncookie { }
  }
  source 0.0.0.0/0
  source-address-translation {
    type automap
  }
  translate-address disabled
  translate-port disabled
}
```

```
ltm profile fastl4 /Common/fl4_swsyncookie {
  app-service none
  defaults-from /Common/fastL4
  hardware-syn-cookie enabled
  pva-dynamic-client-packets 2
  pva-dynamic-server-packets 2
  pva-offload-dynamic enabled
  software-syn-cookie enabled
}
```

```
net route /Common/default {
  gw 10.2.72.66
  network default
}
```

```
-----BIGIP backend as server
net vlan /Common/esnet-1102 {
```

```

interfaces {
    1.1 {
        tag-mode service
        tagged
    }
}
tag 1102
}

net self /Common/10.2.72.66 {
    address 10.2.72.66/16
    allow-service all
    traffic-group /Common/traffic-group-local-only
    vlan /Common/esnet-1102
}

net route /Common/route_to_10.1 {
    gw 10.2.72.12
    network 10.1.0.0/16
}

ltm virtual /Common/vs3_http {
    destination /Common/10.2.72.97:80
    .....SKIP.....
}

ltm virtual /Common/vs98_http {
    destination /Common/10.2.72.98:80
    .....SKIP.....
}

```

4 Ported mTCP/DPDK to KVM VM using with VIRTIO PMD (mTCP/DPDK VM <----> BIGIP VM)

Host hardware : Dell Desktop Optiplex 990

8 core: Intel(R) Core(TM) i7-2600 CPU @ 3.40GHz

16GB RAM

NIC:

03:00:0 Ethernet controller: Broadcom Corporation NetXtreme BCM5720 Gigabit Ethernet PCIe

03:00:1 Ethernet controller: Broadcom Corporation NetXtreme BCM5720 Gigabit Ethernet PCIe

a: run mTCP DPDK app on load VM (generate ~340Mbps):

```
root@dpgk-kvm:/home/dpgk/mtcp# ./apps/example/epwget 10.9.9.6 100000000 -N 1 -c 100000
```

[ALL] connect: 62890, read: 3 MB, write: 6 MB, completes: 46922 (resp_time avg: 9893, max: 2504881 us)

[CPU 0] dpdk0 flows: 100382, RX: 253506(pps) (err: 0), 0.23(Gbps), TX: 384867(pps), 0.34(Gbps)

[ALL] dpdk0 flows: 100382, RX: 253506(pps) (err: 0), 0.23(Gbps), TX: 384867(pps), 0.34(Gbps)

b: BIGIP VM on KVM

Ltm::Virtual Server: vs_http

Status

CMP : enabled
CMP Mode : all-cpus
Destination : 10.9.9.6:80

Traffic ClientSide Ephemeral General

Bits In	233.3G	0	-
Bits Out	184.2G	0	-
Packets In	396.3M	0	-
Packets Out	316.2M	0	-
Current Connections	254.2K	0	-
Maximum Connections	268.3K	0	-

top - 13:57:14 up 1:13, 2 users, load average: 1.96, 1.16, 0.63

Tasks: 411 total, 3 running, 408 sleeping, 0 stopped, 0 zombie

Cpu0 : 31.9%us, 16.9%sy, 0.0%ni, 12.2%id, 0.0%wa, 0.4%hi, 36.6%si, 2.0%st

Cpu1 : 53.4%us, 22.0%sy, 0.0%ni, 21.1%id, 0.0%wa, 0.0%hi, 0.0%si, 3.5%st

Cpu2 : 52.4%us, 20.7%sy, 0.3%ni, 22.0%id, 0.0%wa, 0.0%hi, 0.0%si, 4.5%st

Cpu3 : 57.1%us, 21.1%sy, 0.0%ni, 19.2%id, 0.0%wa, 0.0%hi, 0.0%si, 2.6%st

Mem: 14402956k total, 14159080k used, 243876k free, 4720k buffers

Swap: 1048572k total, 81464k used, 967108k free, 341004k cached

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
13717	root	RT	0	12.1g	127m	102m	R	82.0	0.9	22:23.88	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088
13830	root	RT	0	12.1g	127m	102m	S	81.7	0.9	16:38.70	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088
13829	root	RT	0	12.1g	127m	102m	R	80.4	0.9	17:03.86	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088
13828	root	RT	0	12.1g	127m	102m	S	80.1	0.9	17:04.08	tmm.0 -T 4 --tmid 0 --npus 4 --platform Z100 -m -s 12088

5, Pktgen on SR-IOV

Moved Dell R710 with Intel 82599 10G NIC to ES lab, created two ubuntu VMs with NIC provisioned by Intel 82599 VF, one VM run pktgen syn flooding,

another VM run pktgen dns flooding, this is to simulate both tcp and udp flooding from same 10G NIC port

Dell R710 hypervisor with Intel 82599 SR-IOV

5a, enable SR-IOV in bios and kernel cmdline

root@pktgen:/home/dpdk/mtcp# cat /proc/cmdline

BOOT_IMAGE=/vmlinuz-3.13.0-32-generic root=/dev/mapper/pktgen--vg-root ro intel_iommu=on ixgbe.max_vfs=2

p2p1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP mode DEFAULT group default qlen 1000

link/ether e8:ea:6a:06:1b:1a brd ff:ff:ff:ff:ff:ff

vf 0 MAC 52:54:00:4c:86:ba, spoof checking on, link-state auto

vf 1 MAC 52:54:00:8b:ea:5f, spoof checking on, link-state auto

5b Use the `virsh nodedev-list` command to get the PCI address of the VF you want to assign and its corresponding PF

```
root@pktgen:/home/dpdk/mtcp# virsh nodedev-list | grep 0000_04_
pci_0000_04_00_0
pci_0000_04_00_1
pci_0000_04_10_0
pci_0000_04_10_1
pci_0000_04_10_2
pci_0000_04_10_3
```

5c use `virsh nodedev-dumpxml` to dump the VF info

```
root@pktgen:/home/dpdk/mtcp# virsh nodedev-dumpxml pci_0000_04_10_0
<device>
  <name>pci_0000_04_10_0</name>
  <path>/sys/devices/pci0000:00/0000:00:05.0/0000:04:10.0</path>
  <parent>pci_0000_00_05_0</parent>
  <driver>
    <name>pci-stub</name>
  </driver>
  <capability type='pci'>
    <domain>0</domain>
    <bus>4</bus>
    <slot>16</slot>
    <function>0</function>
    <product id='0x10ed'>82599 Ethernet Controller Virtual Function</product>
    <vendor id='0x8086'>Intel Corporation</vendor>
    <capability type='phys_function'>
      <address domain='0x0000' bus='0x04' slot='0x00' function='0x0'>
    </capability>
    <iommuGroup number='33'>
      <address domain='0x0000' bus='0x04' slot='0x10' function='0x0'>
    </iommuGroup>
  </capability>
</device>
```

The following data is needed for the next step:

```
<domain>0</domain>
```

```
<bus>4</bus>
```

```
<slot>16</slot>
```

```
<function>0</function>
```

Create a temporary XML file (for example `/tmp/vf-interface.xml`) containing the data necessary to add a VF network device to an existing VM Guest. The minimal content of the file needs to look like the following:

```
<interface type='hostdev'>
  <source>
    <address type='pci' domain='0' bus='4' slot='16' function='0'>
  </source>
</interface>
```

5d Specify the data you acquired in the previous step here

In case a device is already attached to the host, it cannot be attached to a guest. To make it available for guests, detach it from the host first:

```
#virsh nodedev-detach pci_0000_04_10_0
```

Last, add the VF interface to an existing VM Guest:

```
#virsh attach-device GUEST /tmp/vf-interface.xml --OPTION
```

GUEST needs to be replaced by the domain name, id or uuid of the VM Guest and `--OPTION` can be one of the following:

--persistent

This option will always add the device to the domain's persistent XML. In addition, if the domain is running, it will be hotplugged.

--config

This option will only affect the persistent XML, even if the domain is running., The device will only show up in the guest on next boot

--live

This option will only affect a running domain. If the domain is inactive, the operation will fail. The device is not persisted in the XML and won't be available in the guest on next boot.

--current

This option affects the current state of the domain. If the domain is inactive, the device is added to the persistent XML and will be available on next boot. If the domain is active, the device is hotplugged but not added to the persistent XML.

5e follow <http://dpdk.org/browse/apps/pktgen-dpdk/tree/README.md> to compile pktgen on VM and binding the ixgbevf to igb_uio

5f run patched pktgen with tcp syn flooding and dns flooding

Note pktgen use core 0 and core 1 for control and screen display, so do not assign logical core 0 and core 1 for packet tx/rx, otherwise, pktgen stops sending packets after a few minutes run

```
root@dppk-sriov:/home/dpdk/pktgen-dpdk/dpdk/examples/pktgen-dpdk# ./app/app/x86_64-native-linuxapp-gcc/pktgen -c 0xf -- -m [2:3].0
```

```
Pktgen> load dnsflood.txt
```

```
Pktgen> start 0
```

```
Pktgen> page range
```

```
Pkten> page main
```

```
MBits/s Rx/Tx :      0/4195      0/4195 <=====~4Gbits/s
```

```
Src/Dst Port :      1234 / 5678
```

```
Pkt Type:VLAN ID:   IPv4 / TCP:0001
```

```
Dst IP Address :      10.6.6.6
```

```
Src IP Address :      10.6.6.128/24
```

```
Dst MAC Address :   00:00:00:00:00:00
```

```
Src MAC Address    52:54:00:4c:86:ba
```

```
-- Pktgen Ver: 2.9.17 (DPDK 16.04.0-rc2) Powered by Intel® DPDK -----
```

```
root@ubuntu-vm:/home/dpdk/pktgen-dpdk/dpdk/examples/pktgen-dpdk# ./app/app/x86_64-native-linuxapp-gcc/pktgen -c 0xf -- -m [2:3].0
```

```
Pktgen> load synflood.txt
```

```
Pktgen> start 0
```

```
Pktgen> page range
```

```
Pkten> page main
```

```
MBits/s Rx/Tx :      0/3739      0/3739 <=====~3Gbits/s
```

6 Pktgen with DPDK net bonding (link aggregation) to BIGIP trunking

NOTE:

6a: don't enable LACP on BIGIP trunk, it appears DPDK net bonding does not work well with LACP with BIGIP trunk

6b: when run Pktgen, make sure the port id is specified correctly to use the net_bonding0 port id, not individual DPDK net bonding member port id

6c: specify DPDK net bonding mode to non mode 4 (LACP) due to 6a

```
root@r710:/home/dpdk/dpdk-upstream/examples/pktgen-dpdk# ./app/app/x86_64-native-linuxapp-gcc/pktgen -c 0xff
--vdev=net_bonding0,mode=0,xmit_policy=134,slave=0000:04:00.1,slave=0000:04:00.0 -- -P -m [0:1-7].2
```

EAL: Initializing pmd_bond for net_bonding0

EAL: Create bonded device net_bonding0 on port 2 in mode 0 on socket 0. <=====here net_bonding0 is port id 2, thus reference this port id in Pktgen -m[0:1-7].2

Pktgen> page stats

```
/          <Real Port Stats Page> Copyright (c) <2010-2016>, Intel Corporation

Port Name          Pkts Rx/Tx    Rx Errors/Missed    Rate Rx/Tx    MAC Address
0-0000:04:00.0:    511/38430488961      0/5067921        0/8085293    E8:EA:6A:06:1B:1B <=====DPDK port 0
1-0000:04:00.1:    511/38430500852      0/5558243        0/8085340    E8:EA:6A:06:1B:1B <=====DPDK port 1
2-net_bonding0:    1022/76860994839      0/10626165       0/16170614    E8:EA:6A:06:1B:1B <=====DPDK net bonding port 2.

-- Pktgen Ver: 3.1.0 (DPDK 17.02.0-rc0) Powered by Intel® DPDK -----
```

References

<http://highscalability.com/blog/2013/5/13/the-secret-to-10-million-concurrent-connections-the-kernel-i.html>

<http://www.dpdk.org/>

<http://www.ndsl.kaist.edu/~kyoungsoo/papers/mtcp.pdf>

<https://github.com/eunyoung14/mtcp>

<https://github.com/emmericp/MoonGen>

<http://git.es.f5net.com/index.cgi/codeshare/tree/vli/mtcp> (git log --author="Vincent Li")

<http://git.es.f5net.com/index.cgi/codeshare/tree/vli/MoonGen>

GS Operations - C10M